

UC Berkeley · CSW182 | [Deep Learning]

Designing, Visualizing and Understanding Deep Neural Networks (2021)

CSW182 (2021) · 课程资料包 @ShowMeAI



视频
中英双语字幕



课件
一键打包下载



笔记
官方笔记翻译



代码
作业项目解析



视频 · B 站 [扫码或点击链接]
<https://www.bilibili.com/video/BV1Ff4y1n7ar>



课件 & 代码 · 博客 [扫码或点击链接]
<http://blog.showmeai.tech/berkeley-csw182>

Berkeley
Q-Learning
计算机视觉

循环神经网络
风格迁移
机器学习基础

可视化
模仿学习
生成模型

梯度策略
元学习
卷积网络

Awesome AI Courses Notes Cheatsheets 是 [ShowMeAI](#) 资料库的分支系列，覆盖最具知名度的 **TOP50+** 门 AI 课程，旨在为读者和学习者提供一整套高品质中文学习笔记和速查表。

点击课程名称，跳转至课程**资料包**页面，**一键下载**课程全部资料！

机器学习	深度学习	自然语言处理	计算机视觉
Stanford · CS229	Stanford · CS230	Stanford · CS224n	Stanford · CS231n
# Awesome AI Courses Notes Cheatsheets · 持续更新中			
知识图谱	图机器学习	深度强化学习	自动驾驶
Stanford · CS520	Stanford · CS224W	UCBerkeley · CS285	MIT · 6.S094



微信公众号

资料下载方式 2：扫码点击**底部菜单栏**
称为 **AI 内容创作者**？回复 [添砖加瓦]

Meta-Learning

Designing, Visualizing and Understanding Deep Neural Networks

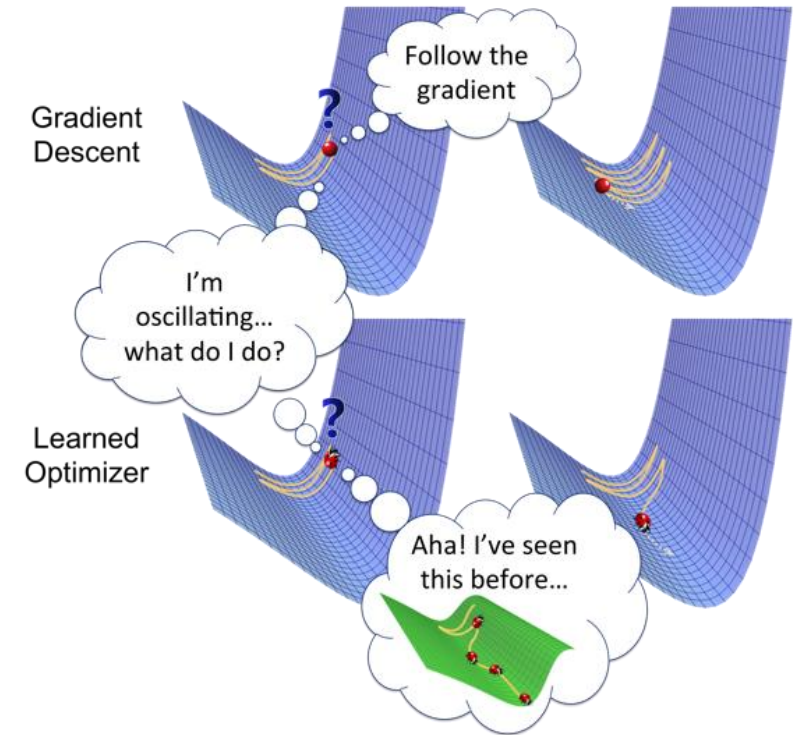
CS W182/282A

Instructor: Sergey Levine
UC Berkeley



What is meta-learning?

- If you've learned 100 tasks already, can you figure out how to *learn* more efficiently?
 - Now having multiple tasks is a huge advantage!
- Meta-learning = *learning to learn*
- In practice, very closely related to multi-task learning
- Many formulations
 - Learning an optimizer
 - Learning an RNN that ingests experience
 - Learning a representation




Why is meta-learning a good idea?

- Deep learning works very well, but requires **large** datasets
- In many cases, we only have a small amount of data available (e.g., some specific computer vision task), but we might have lots of data of a similar type for other tasks (e.g., other object classification tasks)
- How does a *meta-learner* help with this?
 - Use plentiful prior tasks to meta-train a model that can learn a new task quickly with only a few examples
 - Collect a small amount of labeled data for the new task
 - Learn a model on this new dataset that generalizes broadly

Meta-learning with supervised learning



Meta-learning with supervised learning



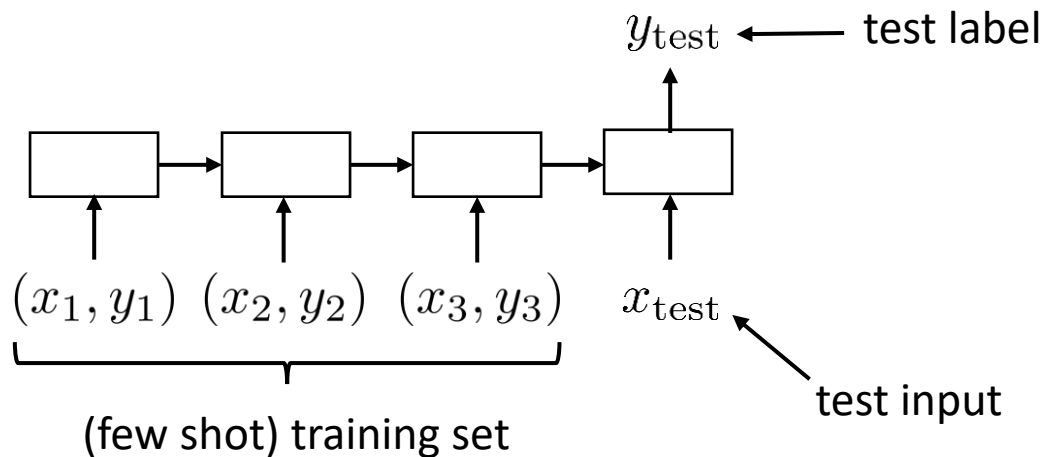
supervised learning: $f(x) \rightarrow y$

input (e.g., image) output (e.g., label)


supervised meta-learning: $f(\mathcal{D}^{\text{tr}}, x) \rightarrow y$

training set

- How to read in training set?
 - Many options, RNNs can work
 - More on this later



What is being “learned”?



supervised meta-learning: $f(\mathcal{D}^{\text{tr}}, x) \rightarrow y$

“Generic” learning:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta, \mathcal{D}^{\text{tr}})$$

$$= f_{\text{learn}}(\mathcal{D}^{\text{tr}})$$

“Generic” meta-learning:

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^n \mathcal{L}(\phi_i, \mathcal{D}_i^{\text{ts}})$$


$$\text{where } \phi_i = f_{\theta}(\mathcal{D}_i^{\text{tr}})$$

What is being “learned”?

“Generic” learning:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta, \mathcal{D}^{\text{tr}})$$


$$= f_{\text{learn}}(\mathcal{D}^{\text{tr}})$$



“Generic” meta-learning:


$$\theta^* = \arg \min_{\theta} \sum_{i=1}^n \mathcal{L}(\phi_i, \mathcal{D}_i^{\text{ts}})$$

$$\text{where } \phi_i = f_{\theta}(\mathcal{D}_i^{\text{tr}})$$

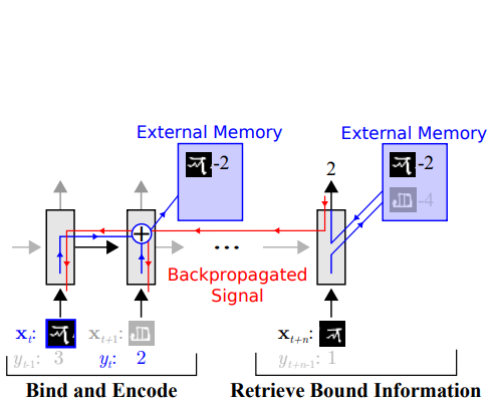


Meta-learning methods


black-box meta-learning



some kind of network that can read in an entire (few-shot) training set

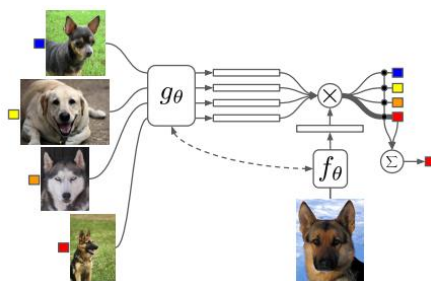


Santoro et al. **Meta-Learning with Memory-Augmented Neural Networks**. 2016.




Mishra et al. **A Simple Neural Attentive Meta-Learner**. 2018.

non-parametric meta-learning




Vinyals et al. **Matching Networks for One Shot Learning**. 2017.



Snell et al. **Prototypical Networks for Few-shot Learning**. 2018.

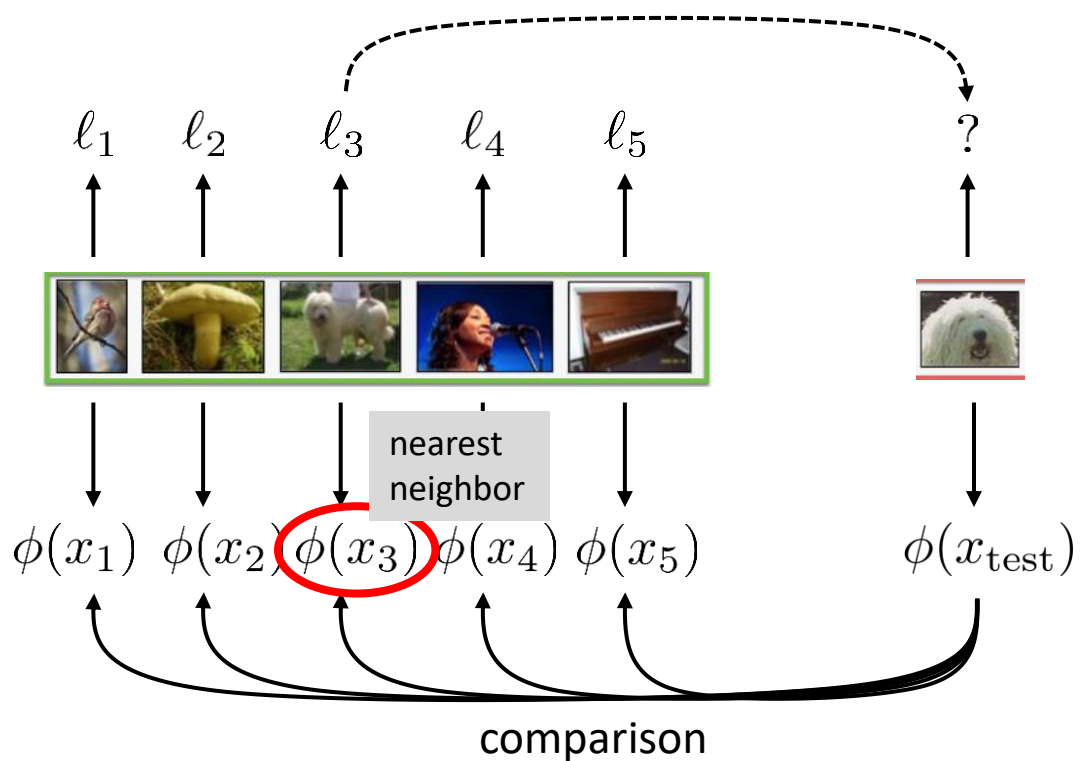
gradient-based meta-learning



Finn et al. **Model-Agnostic Meta-Learning**. 2018.

Non-Parametric & Gradient-Based Meta-Learning

Basic idea



why does this work?

that is, why does the nearest neighbor have the right class?

because we **meta-train** the features so that this produces the right answer!

$$p_{\text{nearest}}(x_k^{\text{tr}} | x_j^{\text{ts}}) \propto \exp(\phi(x_k^{\text{tr}})^T \phi(x_j^{\text{ts}}))$$

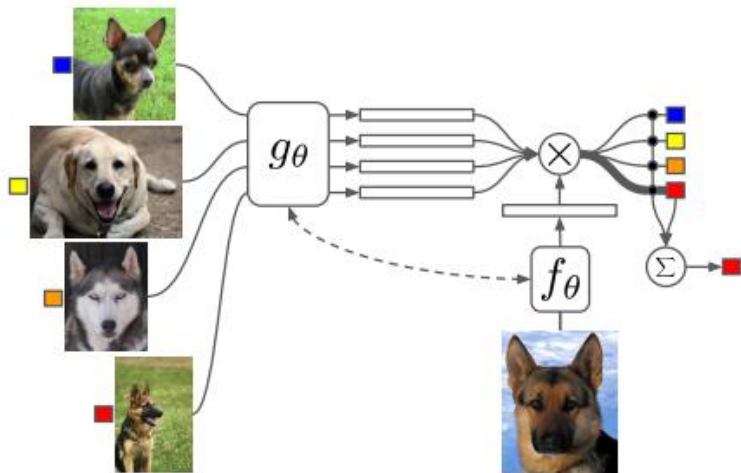
$$p_{\theta}(y_j^{\text{ts}} | x_j^{\text{ts}}, \mathcal{D}_i^{\text{tr}}) = \sum_{k: y_k^{\text{tr}} = y_j^{\text{ts}}} p_{\text{nearest}}(x_k^{\text{tr}} | x_j^{\text{ts}})$$

all training points that have this label

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^n \mathcal{L}(\underbrace{f_{\theta}(\mathcal{D}_i^{\text{tr}})}_{\text{learned (soft) nearest neighbor classifier}}, \mathcal{D}_i^{\text{ts}}) = - \sum_{i=1}^n \sum_{j=1}^m \log p_{\theta}(y_j^{\text{ts}} | x_j^{\text{ts}}, \mathcal{D}_i^{\text{tr}})$$

learned (soft) nearest neighbor classifier

Matching networks



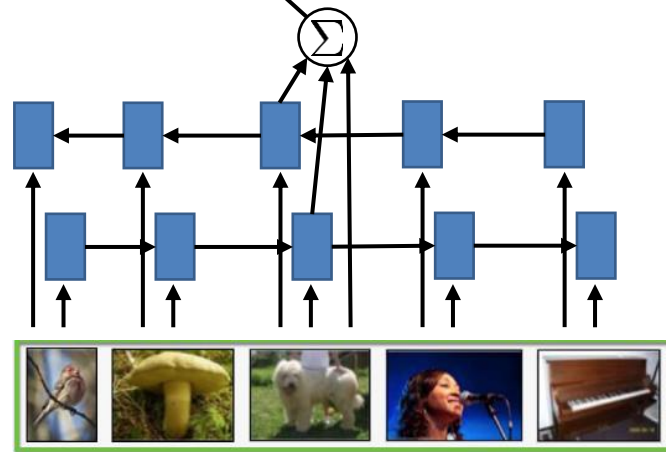
$$p_{\theta}(y_j^{\text{ts}} | x_j^{\text{ts}}, \mathcal{D}_i^{\text{tr}}) = \sum_{k: y_k^{\text{tr}} = y_j^{\text{ts}}} p_{\text{nearest}}(x_k^{\text{tr}} | x_j^{\text{ts}})$$

$$p_{\text{nearest}}(x_k^{\text{tr}} | x_j^{\text{ts}}) \propto \exp(\underbrace{g(x_k^{\text{tr}}, \mathcal{D}_i^{\text{tr}})}_{\text{red}}^T \underbrace{f(x_j^{\text{ts}}, \mathcal{D}_i^{\text{tr}})}_{\text{blue}})$$

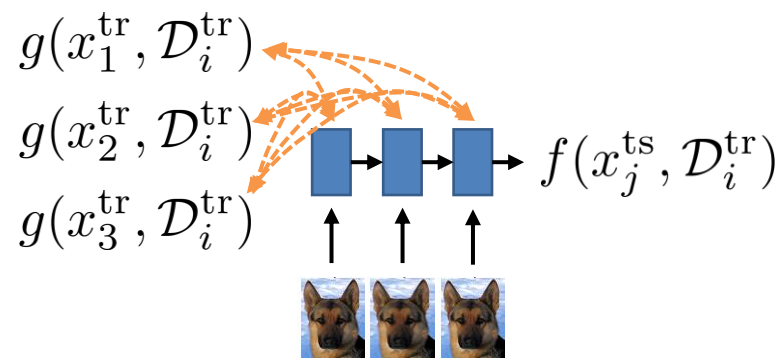
different nets to embed x^{tr} and x^{ts}

both f and g conditioned on entire set $\mathcal{D}_i^{\text{tr}}$


$g(x_k^{\text{tr}}, \mathcal{D}_i^{\text{tr}})$ bidirectional LSTM embedding



$f(x_j^{\text{ts}}, \mathcal{D}_i^{\text{tr}})$ attentional LSTM embedding



Prototypical networks



Two simple ideas compared to matching networks:

1. Instead of “soft nearest neighbor,” construct prototype for each class


$$p_{\theta}(y|x_j^{\text{ts}}, \mathcal{D}_i^{\text{tr}}) \propto \exp(c_y^T f(x_j^{\text{ts}})) \quad c_y = \frac{1}{N_y} \sum_{k:y_k^{\text{tr}}=y} g(x_k^{\text{tr}})$$

2. Get rid of all the complex junk

~~bidirectional LSTM embedding~~

~~attentional LSTM embedding~~

Back to representations...



is pretraining a *type* of meta-learning?

better features = faster learning of new task!

Meta-learning as an optimization problem

$$\theta^* = \arg \min_{\theta} \sum_{i=1}^n \mathcal{L}(\phi_i, \mathcal{D}_i^{\text{ts}})$$

where $\phi_i = f_{\theta}(\mathcal{D}_i^{\text{tr}})$


what if $f_{\theta}(\mathcal{D}_i^{\text{tr}})$ is just a *finetuning* algorithm?

$$f_{\theta}(\mathcal{D}_i^{\text{tr}}) = \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}_i^{\text{tr}})$$

(could take a few gradient steps in general)



This can be trained the same way as any other neural network, by implementing gradient descent as a computation graph and then running backpropagation *through* gradient descent!

MAML in pictures



$$\theta \leftarrow \theta - \beta \sum_i \nabla_{\theta} \mathcal{L}(\theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}_i^{\text{tr}}), \mathcal{D}_i^{\text{ts}})$$

~~$\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta, \mathcal{D}^{\text{tr}})$~~



What did we just do??

supervised learning: $f(x) \rightarrow y$

supervised meta-learning: $f(\mathcal{D}^{\text{tr}}, x) \rightarrow y$

model-agnostic meta-learning: $f_{\text{MAML}}(\mathcal{D}^{\text{tr}}, x) \rightarrow y$

$$f_{\text{MAML}}(\mathcal{D}^{\text{tr}}, x) = f_{\theta'}(x)$$

$$\theta' = \theta - \alpha \sum_{(x,y) \in \mathcal{D}^{\text{tr}}} \nabla_{\theta} \mathcal{L}(f_{\theta}(x), y)$$

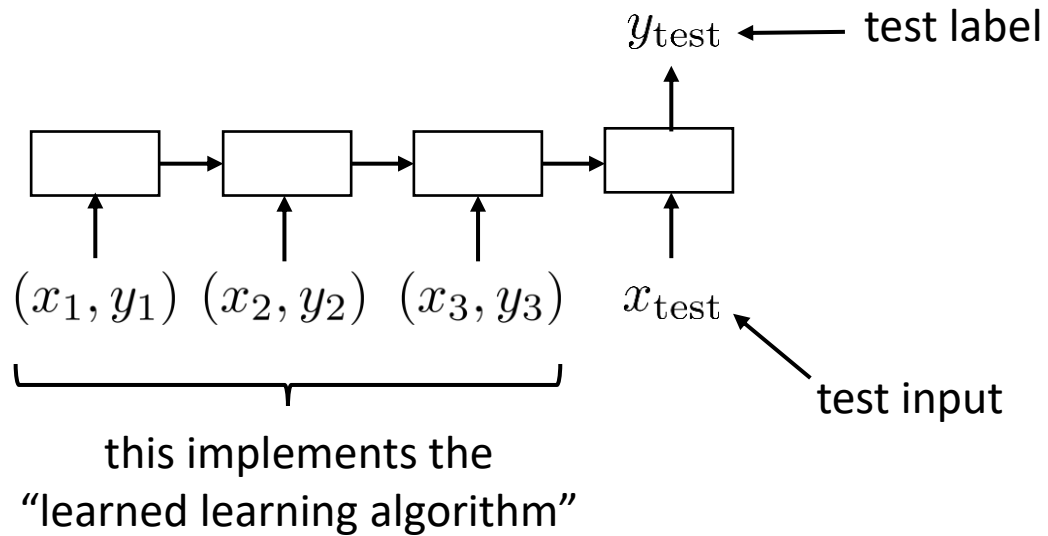
Just another computation graph...

Can implement with any autodiff package (e.g., TensorFlow)

But has favorable inductive bias...


Why does it work?

black-based meta-learning



- Does it converge?
 - Kind of?
- What does it converge to?
 - Who knows...
- What to do if it's not good enough?
 - Nothing...

MAML



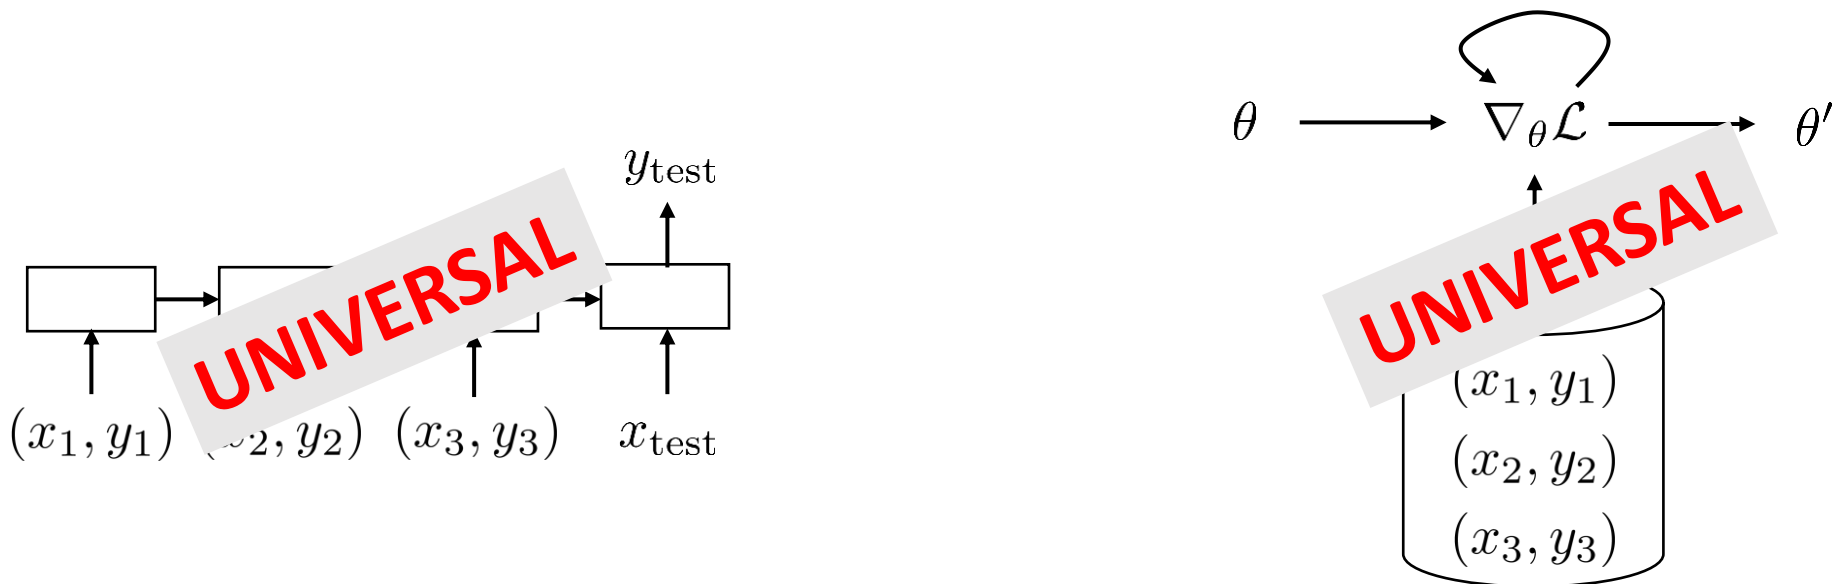
- Does it converge?
 - Yes (it's gradient descent...)
- What does it converge to?
 - A local optimum (it's gradient descent...)
- What to do if it's not good enough?
 - Keep taking gradient steps (it's gradient descent...)

Universality

Did we lose anything?


Universality: meta-learning can learn any “algorithm”

more precisely, can represent any function $f(\mathcal{D}_{\text{train}}, x)$



Summary

black-box meta-learning

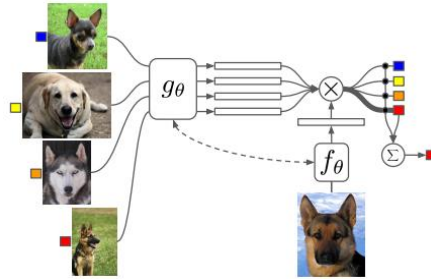


some kind of network that can read in an entire (few-shot) training set

- + conceptually very simple
- + benefits from advances in sequence models (e.g., transformers)

- minimal inductive bias (i.e., *everything* has to be meta-learned)
- hard to scale to “medium” shot (we get long “sequences”)


non-parametric meta-learning



Vinyals et al. **Matching Networks for One Shot Learning**. 2017.

- + can work very well by combining some inductive bias with easy end-to-end optimization
- restricted to classification, hard to extend to other settings like regression or reinforcement learning
- somewhat specialized architectures

gradient-based meta-learning



Finn et al. **Model-Agnostic Meta-Learning**. 2018.

- + easy to apply to any architecture or loss function (inc. RL, regression)
- + good generalization to out-of-domain tasks
- meta-training optimization problem is harder, requires more tuning
- requires second derivatives

Meta-Reinforcement Learning

The meta reinforcement learning problem

“Generic” learning:

$$\begin{aligned}\theta^* &= \arg \min_{\theta} \mathcal{L}(\theta, \mathcal{D}^{\text{tr}}) \\ &= f_{\text{learn}}(\mathcal{D}^{\text{tr}})\end{aligned}$$

Reinforcement learning:

$$\begin{aligned}\theta^* &= \arg \max_{\theta} E_{\pi_{\theta}(\tau)}[R(\tau)] \\ &= f_{\text{RL}}(\mathcal{M})\end{aligned}\quad \mathcal{M} = \{\mathcal{S}, \mathcal{A}, \mathcal{P}, r\}$$

↖
MDP

“Generic” meta-learning:

$$\begin{aligned}\theta^* &= \arg \min_{\theta} \sum_{i=1}^n \mathcal{L}(\phi_i, \mathcal{D}_i^{\text{ts}}) \\ &\text{where } \phi_i = f_{\theta}(\mathcal{D}_i^{\text{tr}})\end{aligned}$$

Meta-reinforcement learning:

$$\begin{aligned}\theta^* &= \arg \max_{\theta} \sum_{i=1}^n E_{\pi_{\phi_i}(\tau)}[R(\tau)] \\ &\text{where } \phi_i = f_{\theta}(\mathcal{M}_i)\end{aligned}$$

↖
MDP for task i

The meta reinforcement learning problem

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^n E_{\pi_{\phi_i}(\tau)} [R(\tau)]$$

where $\phi_i = f_{\theta}(\mathcal{M}_i)$

assumption: $\mathcal{M}_i \sim p(\mathcal{M})$


meta test-time:

sample $\mathcal{M}_{\text{test}} \sim p(\mathcal{M})$, get $\phi_i = f_{\theta}(\mathcal{M}_{\text{test}})$

$\{\mathcal{M}_1, \dots, \mathcal{M}_n\}$

↑
meta-training MDPs

Some examples:



Meta-RL with recurrent policies


$$\theta^* = \arg \max_{\theta} \sum_{i=1}^n E_{\pi_{\phi_i}(\tau)}[R(\tau)]$$

where $\phi_i = f_{\theta}(\mathcal{M}_i)$

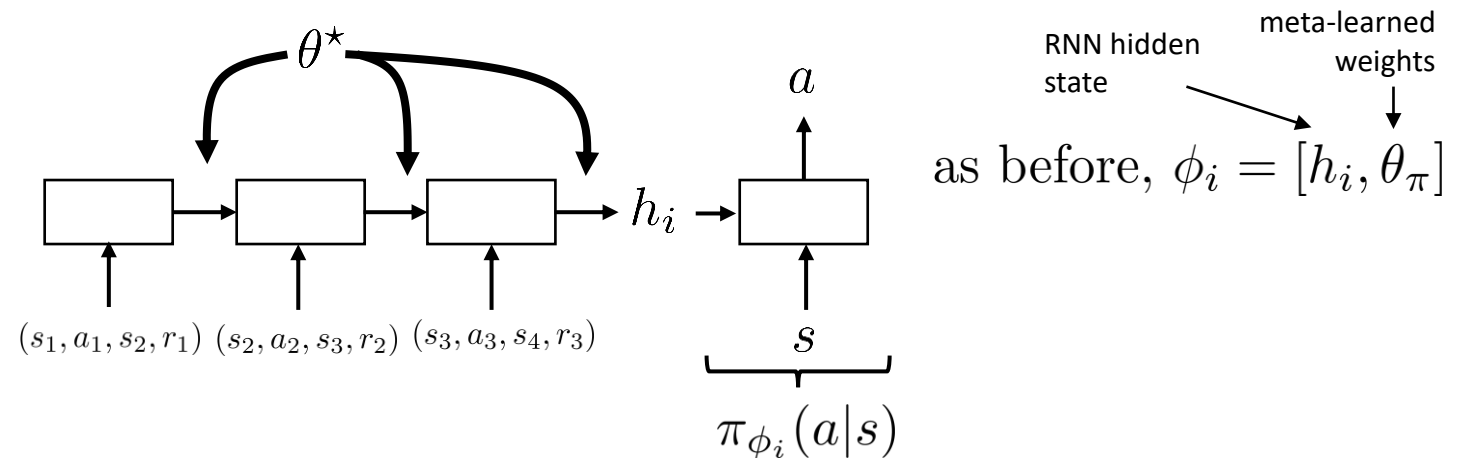
main question: how to implement $f_{\theta}(\mathcal{M}_i)$?

what should $f_{\theta}(\mathcal{M}_i)$ do?

1. improve policy with experience from \mathcal{M}_i
 $\{(s_1, a_1, s_2, r_1), \dots, (s_T, a_T, s_{T+1}, r_T)\}$
2. (new in RL): choose how to interact, i.e. choose a_t
 meta-RL must also *choose* how to *explore*!




use (s_t, a_t, s_{t+1}, r_t) to improve π_{θ}



Meta-RL with recurrent policies

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^n E_{\pi_{\phi_i}(\tau)}[R(\tau)]$$



where $\phi_i = f_{\theta}(\mathcal{M}_i)$




so... we just train an RNN policy?

yes!



crucially, RNN hidden state is **not** reset between episodes!



Why recurrent policies *learn to explore*



1. improve policy with experience from \mathcal{M}_i
 $\{(s_1, a_1, s_2, r_1), \dots, (s_T, a_T, s_{T+1}, r_T)\}$
2. (new in RL): choose how to interact, i.e. choose a_t
 meta-RL must also *choose* how to *explore*!






$$\theta^* = \arg \max_{\theta} E_{\pi_{\theta}} \left[\sum_{t=0}^T r(s_t, a_t) \right]$$

optimizing total reward over the entire **meta**-episode with RNN policy **automatically** learns to explore!


Meta-RL with recurrent policies

$$\theta^* = \arg \max_{\theta} \sum_{i=1}^n E_{\pi_{\phi_i}(\tau)} [R(\tau)]$$


where $\phi_i = f_{\theta}(\mathcal{M}_i)$



(a) Labryinth I-maze



(b) Illustrative Episode



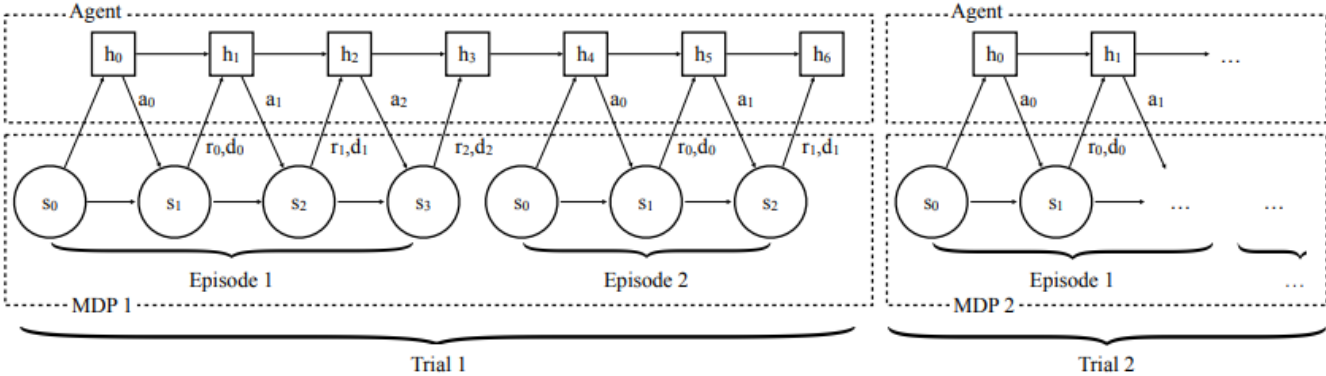
(a) Good behavior, 1st episode (b) Good behavior, 2nd episode (c) Bad behavior, 1st episode (d) Bad behavior, 2nd episode

Heess, Hunt, Lillicrap, Silver. **Memory-based control with recurrent neural networks**. 2015.

Wang, Kurth-Nelson, Tirumala, Soyer, Leibo, Munos, Blundell, Kumaran, Botvinick. **Learning to Reinforcement Learning**. 2016.


Duan, Schulman, Chen, Bartlett, Sutskever, Abbeel. **RL2: Fast Reinforcement Learning via Slow Reinforcement Learning**. 2016.

Architectures for meta-RL




standard RNN (LSTM) architecture

Duan, Schulman, Chen, Bartlett, Sutskever, Abbeel. **RL2: Fast Reinforcement Learning via Slow Reinforcement Learning**. 2016.



attention + temporal convolution

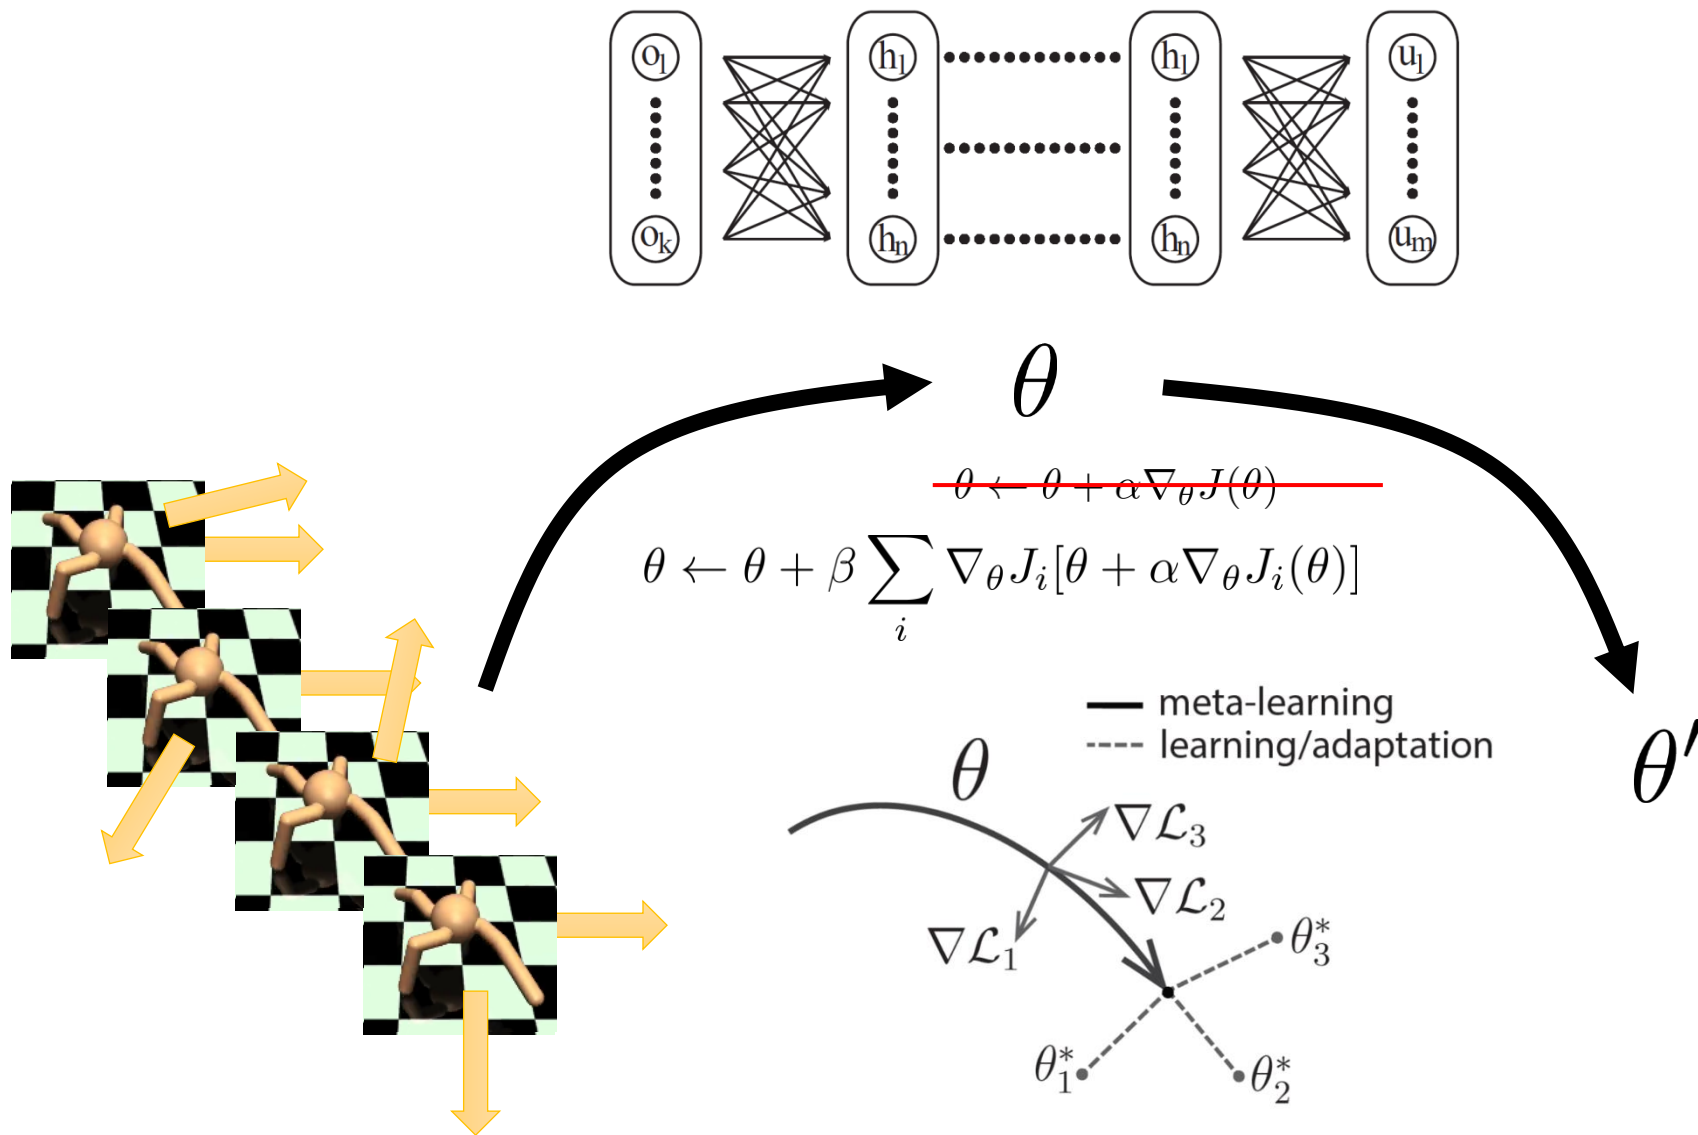
Mishra, Rohaninejad, Chen, Abbeel. **A Simple Neural Attentive Meta-Learner**.



parallel permutation-invariant context encoder



Rakelly*, Zhou*, Quillen, Finn, Levine. **Efficient Off-Policy Meta-Reinforcement learning via Probabilistic Context Variables**.

MAML for RL





MAML for RL videos



after MAML training



after 1 gradient step
(forward reward)



after 1 gradient step
(backward reward)



UC Berkeley · CSW182 | [Deep Learning]

Designing, Visualizing and Understanding Deep Neural Networks (2021)

CSW182 (2021) · 课程资料包 @ShowMeAI



视频
中英双语字幕



课件
一键打包下载



笔记
官方笔记翻译



代码
作业项目解析



视频 · B 站 [扫码或点击链接]
<https://www.bilibili.com/video/BV1Ff4y1n7ar>



课件 & 代码 · 博客 [扫码或点击链接]
<http://blog.showmeai.tech/berkeley-csw182>

Berkeley
Q-Learning
计算机视觉

循环神经网络
风格迁移
机器学习基础

可视化
模仿学习
生成模型

梯度策略
元学习
卷积网络

Awesome AI Courses Notes Cheatsheets 是 [ShowMeAI](#) 资料库的分支系列，覆盖最具知名度的 **TOP50+** 门 AI 课程，旨在为读者和学习者提供一整套高品质中文学习笔记和速查表。

点击课程名称，跳转至课程**资料包**页面，**一键下载**课程全部资料！

机器学习	深度学习	自然语言处理	计算机视觉
Stanford · CS229	Stanford · CS230	Stanford · CS224n	Stanford · CS231n
# Awesome AI Courses Notes Cheatsheets · 持续更新中			
知识图谱	图机器学习	深度强化学习	自动驾驶
Stanford · CS520	Stanford · CS224W	UCBerkeley · CS285	MIT · 6.S094



微信公众号

资料下载方式 2：扫码点击**底部菜单栏**
称为 **AI 内容创作者**？回复 [添砖加瓦]